

# FIT 3080: Intelligent Systems

## Markov Decision Processes (MDPs)

Gholamreza Haffari – Monash University

Many slides over the course adapted from Stuart Russell,  
Andrew Moore, or Dan Klein

# Announcements

---


- Online Reading:

- Reinforcement Learning: An Introduction, by Richard Sutton and Andrew Barto, MIT Press
- Chapter 3 and Chapter 4
- Accessible from:  
<http://webdocs.cs.ualberta.ca/~sutton/book/ebook/the-book.html>
- Different treatment and notation than the R&N book, beware!
- Lecture version is the standard for this class



- R&N book:

- Sections 17.1-3



# Non-Deterministic Search

---

- How do you plan when your actions might fail?



# Outline

---

- MDPs
  - Definition
  - MDPs as Search Trees
  - Utility of a Sequence of Actions
- Solving MDPs
  - Bellman Equation
  - Value Iteration
  - Policy Iteration

# Outline


---

- MDPs
  - Definition
  - MDPs as Search Trees
  - Utility of a Sequence of Actions
- Solving MDPs
  - Bellman Equation
  - Value Iteration
  - Policy Iteration

# Grid World


- The agent lives in a grid

- Agent has four actions: North, South, East, West
- Walls block the agent's path
- If outcome of actions are deterministic, we can use standard planning/search algorithms to go from the initial state to the favorable goal state



- The agent's actions do not always go as planned:


- 80% of the time, the action North takes the agent North (if there is no wall there)
- 10% of the time, North takes the agent West; 10% East
- If there is a wall in the direction the agent would have been taken, the agent stays put




# Action Results

---

Deterministic Grid World




Non-Deterministic Grid World



# Grid World

---

- Small “living” reward each step
- Big rewards come at the end
- **Goal: maximize sum of rewards\***





- **Question: What is the best action to take from each state to maximize sum of rewards?**
- **This week lectures:**
  - Firstly, formalize the problem more generally as a Markov Decision Process (MDP)
  - Secondly, present algorithms to solve MDPs (to solve stochastic planning/game problems)



# Markov Decision Processes

- An MDP is defined by:
  - A **set of states**  $s \in S$
  - A **set of actions**  $a \in A$
  - A **transition function**  $T(s, a, s')$ 
    - Prob that  $a$  from  $s$  leads to  $s'$
    - i.e.,  $P(s' | s, a)$
    - Also called the model
  - A **reward function**  $R(s, a, s')$ 
    - Sometimes just  $R(s)$  or  $R(s')$
  - A **start state** (or distribution)
  - Maybe a **terminal state**
- MDPs are a family of non-deterministic search problems
  - Reinforcement learning: MDPs where we don't know the transition or reward functions



# What is Markov about MDPs?

---

- Andrey Markov (1856-1922)
- “Markov” generally means that given the present state, the future and the past are independent
- For Markov decision processes, “Markov” means:



$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1}, \dots, S_0 = s_0)$$


=

$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t)$$

# Solving MDPs


- In deterministic single-agent search problems, want an optimal **plan**, or sequence of actions, from start to a goal
- In an MDP, we want an optimal **policy**  $\pi^*: S \rightarrow A$ 
  - A policy  $\pi$  gives an action for each state
  - An optimal policy maximizes expected utility if followed
  - Defines a reflex agent

Optimal policy when  
 $R(s, a, s') = -0.03$  for  
all non-terminals  $s$




# Example Optimal Policies


---




$$R(s) = -0.01$$



$$R(s) = -0.03$$



$$R(s) = -0.4$$




$$R(s) = -2.0$$

# Example: High-Low

---

- Rules:

- Three card types: 2, 3, 4
  - Infinite deck, twice as many 2's
- Start with 3 showing
- After each card, you say “high” or “low”
  - New card is flipped
  - If you're right, you win the points shown on the new card
  - Ties are no-ops
  - If you're wrong, game ends



- How is this different from the “chance” games in the last lecture?


- #1: get rewards as you go
- #2: you might play forever

You can patch expectimax to deal with #1 but not #2

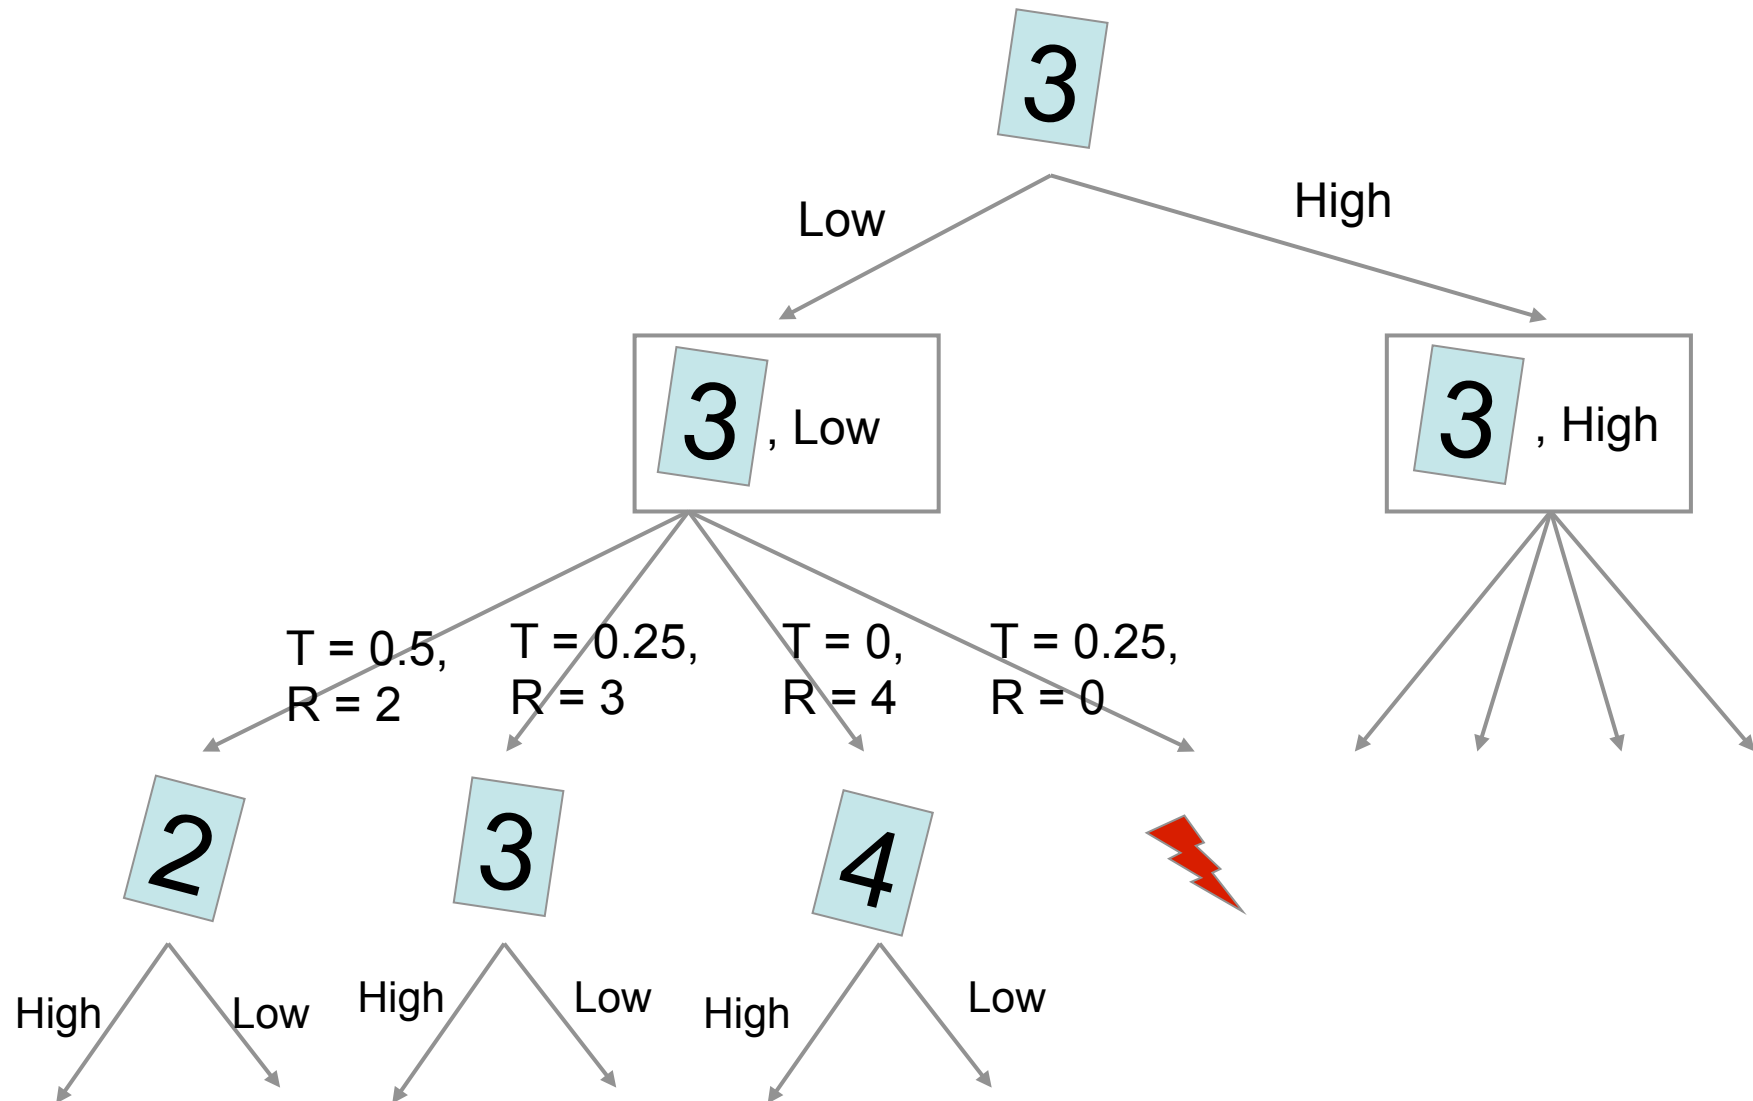
# High-Low as an MDP

---

- States: 2, 3, 4, done
- Actions: High, Low
- Model:  $T(s, a, s')$ :
  - $P(s' = 4 \mid 4, \text{Low}) = 1/4$
  - $P(s' = 3 \mid 4, \text{Low}) = 1/4$
  - $P(s' = 2 \mid 4, \text{Low}) = 1/2$
  - $P(s' = \text{done} \mid 4, \text{Low}) = 0$
  - $P(s' = 4 \mid 4, \text{High}) = 1/4$
  - $P(s' = 3 \mid 4, \text{High}) = 0$
  - $P(s' = 2 \mid 4, \text{High}) = 0$
  - $P(s' = \text{done} \mid 4, \text{High}) = 3/4$
  - ...
- Rewards:  $R(s, a, s')$ :
  - Number shown on  $s'$  if  $s \neq s'$
  - 0 otherwise



# Example: High-Low



# Outline


---

- MDPs
  - Definition
  - MDPs as Search Trees
  - Utility of a Sequence of Actions
- Solving MDPs
  - Bellman Equation
  - Value Iteration
  - Policy Iteration



# MDP Search Trees

- Each MDP state gives an “expectimax” search tree



# Outline

---

- MDPs
  - Definition
  - MDPs as Search Trees
  - Utility of a Sequence of Actions
- Solving MDPs
  - Bellman Equation
  - Value Iteration
  - Policy Iteration

# Utilities of Sequences

---

- In order to formalize optimality of a policy, need to understand utilities of sequences of rewards
- Typically consider **stationary preferences**:

$$\begin{aligned} [r, r_0, r_1, r_2, \dots] &\succ [r, r'_0, r'_1, r'_2, \dots] \\ &\Leftrightarrow \\ [r_0, r_1, r_2, \dots] &\succ [r'_0, r'_1, r'_2, \dots] \end{aligned}$$

- **Theorem: only two ways to define stationary utilities**
  - Additive utility:

$$U([r_0, r_1, r_2, \dots]) = r_0 + r_1 + r_2 + \dots$$

- Discounted utility:

$$U([r_0, r_1, r_2, \dots]) = r_0 + \gamma r_1 + \gamma^2 r_2 + \dots$$

# Infinite Utilities?!

- Problem: infinite state sequences have infinite rewards


- Solutions:

- Finite horizon:

- Terminate episodes after a fixed T steps (e.g. life)
    - Gives nonstationary policies ( $\pi$  depends on time left)

- Absorbing state: guarantee that for every policy, a terminal state will eventually be reached (like “done” for High-Low)

- Discounting: for  $0 < \gamma < 1$




$$U([r_0, \dots, r_\infty]) = \sum_{t=0}^{\infty} \gamma^t r_t \leq R_{\max} / (1 - \gamma)$$

- Smaller  $\gamma$  means smaller “horizon” – shorter term focus

# Discounting

- Typically discount rewards by  $\gamma < 1$  each time step
  - Sooner rewards have higher utility than later rewards
  - Also helps the algorithms converge



# Outline

---


- MDPs
  - Definition
  - MDPs as Search Trees
  - Utility of a Sequence of Actions
- Solving MDPs
  - Bellman Equation
  - Value Iteration
  - Policy Iteration

# Recap: Defining MDPs

---

- Markov decision processes:

- States  $S$
- Start state  $s_0$
- Actions  $A$
- Transitions  $P(s' | s, a)$  (or  $T(s, a, s')$ )
- Rewards  $R(s, a, s')$  (and discount  $\gamma$ )




- MDP quantities so far:

- Policy = Choice of action for each state
- Utility (or return) = sum of discounted rewards

# Optimal Utilities

- Define the value of a state  $s$ :  
 $V^*(s)$  = expected utility starting in  $s$  and acting optimally
- Define the value of a q-state  $(s,a)$ :  
 $Q^*(s,a)$  = expected utility starting in  $s$ , taking action  $a$  and thereafter acting optimally
- Define the optimal policy:  
 $\pi^*(s)$  = optimal action from state  $s$



3	0.812	0.868	0.912	$+1$
2	0.762		0.660	$-1$
1	0.705	0.655	0.611	0.388
	1	2	3	4

3	→	→	→	$+1$
2	↑		↑	$-1$
1	↑	←	←	←
	1	2	3	4



# Solving MDPs


---

- We want to find the optimal policy  $\pi^*$
- How to do that?

# The Bellman Equations

- Definition of “optimal utility” leads to a simple one-step lookahead relationship amongst optimal utility values:

Optimal rewards = maximize over first action and then follow optimal policy



- Formally:


$$V^*(s) = \max_a Q^*(s, a)$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

# Value Estimates

- Calculate estimates  $V_k^*(s)$ 
  - Not the optimal value of  $s$ !
  - The optimal value considering only next  $k$  time steps ( $k$  rewards)
  - As  $k \rightarrow \infty$ , it approaches the optimal value
- Almost solution: recursion (i.e. expectimax)
- Correct solution: dynamic programming



# Outline

---

- MDPs
  - Definition
  - MDPs as Search Trees
  - Utility of a Sequence of Actions
- Solving MDPs
  - Bellman Equation
  - Value Iteration
  - Policy Iteration

# Value Iteration

---

- Idea:


- Start with  $V_0^*(s) = 0$ , which we know is right (why?)
- Given  $V_i^*$ , calculate the values for all states for depth  $i+1$ :

$$V_{i+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i(s')]$$

- This is called a **value update** or **Bellman update**
  - Repeat until convergence
- Theorem: will converge to unique optimal values
    - Basic idea: approximations get refined towards optimal values
    - Policy may converge long before values do

Example:  $\gamma=0.9$ , living  
reward=0, noise=0.2

# Example: Bellman Updates



$$V_{i+1}(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i(s')]$$


$$V_2(\langle 3, 3 \rangle) = \sum_{s'} T(\langle 3, 3 \rangle, \text{right}, s') [R(\langle 3, 3 \rangle) + 0.9 V_1(s')]$$

max happens for  
 $a=\text{right}$ , other  
actions not shown

$$= 0.9 [0.8 \cdot 1 + 0.1 \cdot 0 + 0.1 \cdot 0]$$

# Example: Value Iteration

---



- Information propagates outward from terminal states and eventually all states have correct value estimates

# Convergence\*

---

- Define the max-norm:  $\|U\| = \max_s |U(s)|$

- Theorem: For any two approximations U and V

$$\|U^{t+1} - V^{t+1}\| \leq \gamma \|U^t - V^t\|$$

- I.e. any distinct approximations must get closer to each other, so, in particular, any approximation must get closer to the true U and value iteration converges to a unique, stable, optimal solution
- Theorem:  
 $\|U^{t+1} - U^t\| < \epsilon, \Rightarrow \|U^{t+1} - U\| < 2\epsilon\gamma/(1 - \gamma)$

- I.e. once the change in our approximation is small, it must also be close to correct



# Practice: Computing Actions

---

- Which action should we chose from state  $s$ :
  - Given optimal values  $V$ ?

$$\arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- Given optimal q-values  $Q$ ?

$$\arg \max_a Q^*(s, a)$$

- Lesson: actions are easier to select from  $Q$ 's!

# Aside: Q-Value Iteration

---

- Value iteration: find successive approx optimal values
  - Start with  $V_0^*(s) = 0$ , which we know is right (why?)
  - Given  $V_i^*$ , calculate the values for all states for depth  $i+1$ :

$$V_{i+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i(s')]$$

- But Q-values are more useful!
  - Start with  $Q_0^*(s, a) = 0$ , which we know is right (why?)
  - Given  $Q_i^*$ , calculate the q-values for all q-states for depth  $i+1$ :

$$Q_{i+1}(s, a) \leftarrow \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma \max_{a'} Q_i(s', a')]$$

# Outline

---

- MDPs
  - Definition
  - MDPs as Search Trees
  - Utility of a Sequence of Actions
- Solving MDPs
  - Bellman Equation
  - Value Iteration
  - Policy Iteration


# Utilities for Fixed Policies

- Another basic operation: compute the utility of a state  $s$  under a fix (general non-optimal) policy
- Define the utility of a state  $s$ , under a fixed policy  $\pi$ :

$V^\pi(s)$  = expected total discounted rewards (return) starting in  $s$  and following  $\pi$

- Recursive relation (one-step look-ahead / Bellman equation):

$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$



# Policy Evaluation

---

- How do we calculate the  $V$ 's for a fixed policy?
- Idea one: modify Bellman updates

$$V_0^\pi(s) = 0$$

$$V_{i+1}^\pi(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_i^\pi(s')]$$

- Idea two: it's just a linear system, solve with Matlab (or whatever)

# Policy Iteration

---

- Problem with value iteration:
  - Considering all actions each iteration is slow: takes  $|A|$  times longer than policy evaluation
  - But policy doesn't change each iteration, time wasted
- Alternative to value iteration:
  - **Step 1: Policy evaluation:** calculate utilities for a fixed policy (not optimal utilities!) until convergence (fast)
  - **Step 2: Policy improvement:** update policy using one-step lookahead with resulting converged (but not optimal!) utilities (slow but infrequent)
  - Repeat steps until policy converges
- This is **policy iteration**
  - It's still optimal!
  - Can converge faster under some conditions

# Policy Iteration

---

- **Policy evaluation:** with fixed current policy  $\pi$ , find values with simplified Bellman updates
  - Iterate until values converge

$$V_{i+1}^{\pi_k}(s) \leftarrow \sum_{s'} T(s, \pi_k(s), s') [R(s, \pi_k(s), s') + \gamma V_i^{\pi_k}(s')]$$

- **Policy improvement:** with fixed utilities, find the best action according to one-step look-ahead

$$\pi_{k+1}(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^{\pi_k}(s')]$$

# Policy Iteration Complexity

---

- Problem size:
  - $|A|$  actions and  $|S|$  states
- Each Iteration
  - Computation:  $O(|S|^3 + |A|.|S|^2)$
  - Space:  $O(|S|)$
- Number of iterations
  - Unknown, but can be faster in practice
  - Convergence is guaranteed



# Comparison

---

- In value iteration:
  - Every pass (or “backup”) updates both utilities (explicitly, based on current utilities) and policy (possibly implicitly, based on current policy)
- In policy iteration:
  - Several passes to update utilities with frozen policy
  - Occasional passes to update policies
- Hybrid approaches (asynchronous policy iteration):
  - Any sequences of partial updates to either policy entries or utilities will converge if every state is visited infinitely often