

# FIT 3080: Intelligent Systems

## Expectimax and Reinforcement Learning

Gholamreza Haffari – Monash University

Many slides over the course adapted from Stuart Russell,  
Andrew Moore, or Dan Klein

# Announcements

---


- Online Reading:

- Reinforcement Learning: An Introduction, by Richard Sutton and Andrew Barto, MIT Press
- Chapter 3 and Chapter 4
- Accessible from:  
<http://webdocs.cs.ualberta.ca/~sutton/book/ebook/the-book.html>
- Different treatment and notation than the R&N book, beware!
- Lecture version is the standard for this class



- R&N book:

- Section 5.5
- Sections 17.1-3



# Outline

---

- Expectimax Search
- Reinforcement Learning (RL)
- Passive Learning in RL
  - Model-based
  - Model-free
    - Direct Estimation
    - Temporal Difference
- Active Learning in RL
  - Q-Learning


# Outline

---

- Expectimax Search
- Reinforcement Learning (RL)
- Passive Learning in RL
  - Model-based
  - Model-free
    - Direct Estimation
    - Temporal Difference
- Active Learning in RL
  - Q-Learning



# Deterministic Games

- Deterministic, zero-sum two-player games:
  - Tic-tac-toe, chess, checkers
  - One player maximizes result
  - The other minimizes result
- Minimax search:
  - A state-space search tree
  - Players alternate turns
  - Each node has a **minimax value**: best achievable **utility** against a rational adversary



# Stochastic/Non-Deterministic Games

- Stochastic games:
  - Backgammon, Solitaire, Minesweeper, ...
- Result of an action can be uncertain
  - eg in Backgammon, before rolling the dice, we don't know what's the outcome
- Can we approach it as search in a state space?
  - What's the utility of an action with uncertain outcomes?




# Utility of an Uncertain Action?

- For uncertain actions, consider the **expected utility**:

$$\text{Utility}(\text{action}) = \sum P(\text{state}|\text{action}) * \text{Utility}(\text{state})$$

- Example:**


- I want to go from home to the airport
- I can take the freeway (action)
- Outcome of take freeway is uncertain:
  - State Arrive Early: (.6, +50)
  - State Arrive Late: (.4, -200)
  - Expected Utility =  $.6 * 50 + .4 * (-200) = -50$



# Expectimax Search Trees

(vs mini-max search trees)

- Can do **expectimax search**
  - Chance nodes, like min nodes, except the outcome is uncertain
  - Calculate **expected utilities**
  - Max nodes as in minimax search
  - Chance nodes take average (expectation) of value of children
- More formally, we have seen how to formalize the underlying problem as a **Markov Decision Process**





# Expectimax Pseudocode

```
def value(s)
```

```
    if s is a max node return maxValue(s)
```

```
    if s is an exp node return expValue(s)
```

```
    if s is a terminal node return evaluation(s)
```

```
def maxValue(s)
```

```
    values = [value(s') for s' in successors(s)]
```



```
    return max(values)
```

```
def expValue(s)
```

```
    values = [value(s') for s' in successors(s)]
```


```
    weights = [probability(s, s') for s' in successors(s)]
```

```
    return expectation(values, weights)
```



# Mixed Layer Types

- E.g. Backgammon
- Expectiminimax
  - Environment is an extra player that moves after each agent
  - Chance nodes take expectations, otherwise like minimax



```
if state is a MAX node then
    return the highest EXPECTIMINIMAX-VALUE of SUCCESSORS(state)
if state is a MIN node then
    return the lowest EXPECTIMINIMAX-VALUE of SUCCESSORS(state)
if state is a chance node then
    return average of EXPECTIMINIMAX-VALUE of SUCCESSORS(state)
```

# Outline

---


- Expectimax Search
- Reinforcement Learning (RL)
- Passive Learning in RL
  - Model-based
  - Model-free
    - Direct Estimation
    - Temporal Difference
- Active Learning in RL
  - Q-Learning

# Recap: MDPs

---

- Markov decision processes:

- States  $S$
- Actions  $A$
- Transitions  $P(s' | s, a)$  (or  $T(s, a, s')$ )
- Rewards  $R(s, a, s')$  (and discount  $\gamma$ )
- Start state  $s_0$  (or distribution  $P_0$ )



- Quantities:


- Policy = map of states to actions
- Utility = sum of discounted rewards
- Values = expected future utility from a state
- Q-Values: expected future utility from a q-state

# Reinforcement Learning

---

- Basic idea:

- Receive feedback in the form of **rewards**
- Agent's utility is defined by the reward function
- Must learn to act so as to **maximize expected rewards**



# Reinforcement Learning

---

- Reinforcement learning:
  - Still assume an MDP:
    - A set of states  $s \in S$
    - A set of actions (per state)  $A$
    - A model  $T(s,a,s')$
    - A reward function  $R(s,a,s')$
  - Still looking for a policy  $\pi(s)$
- New twist: don't know  $T$  or  $R$ 
  - I.e. don't know which states are good or what the actions do
  - Must actually try actions and states out to learn



# Example: Animal Learning

---

- RL studied experimentally for more than 60 years in psychology
  - Rewards: food, pain, hunger, drugs, etc.
  - Mechanisms and sophistications debated
- Example: foraging
  - Bees learn near-optimal foraging plan in field of artificial flowers with controlled nectar supplies
  - Bees have a direct neural connection from nectar intake measurement to motor planning area

# Example: Backgammon

- Reward only for win / loss in terminal states, zero otherwise
  - **TD-Gammon** learns a function approximation to  $V(s)$  using **a neural network**
  - Combined with depth 3 search, one of the top 3 players in the world
- 
- You could imagine training Pacman this way ...
  - But it's tricky!





# Key Ideas for Learning

---

- Online vs. Batch

- Learn while exploring the world, or learn from fixed batch of data

- Active vs. Passive

- Does the learner actively choose actions to gather experience? Or, is a fixed policy provided?

- Model learning vs. Model free

- Do we estimate  $T(s,a,s')$  and  $R(s,a,s')$  , or just learn values/policy directly

# Outline


---

- Expectimax Search
- Reinforcement Learning (RL)
- **Passive Learning in RL**
  - Model-based
  - Model-free
    - Direct Estimation
    - Temporal Difference
- Active Learning in RL
  - Q-Learning

# Passive Learning

- Simplified task

- You don't know the transitions  $T(s,a,s')$
- You don't know the rewards  $R(s,a,s')$
- You are given a policy  $\pi(s)$
- Goal: learn the state values
- ... what value iteration did!



- In this case:

- Learner has no choice about what actions to take
- Just execute the policy and learn from experience
- We'll get to the active case soon

# Detour: Sampling Expectations

---

- What is the average height of people in Monash?
- Method: measure their heights, add them up, and divide by  $N$



# Detour: Sampling Expectations

---

- Want to compute an expectation weighted by  $P(x)$ :

$$E[f(x)] = \sum_x P(x) f(x)$$

- Model-based:** estimate  $P(x)$  from samples, compute expectation

$$\begin{aligned} x_i &\sim P(x) \\ \hat{P}(x) &= \text{count}(x)/k \end{aligned} \qquad E[f(x)] \approx \sum_x \hat{P}(x) f(x)$$

- Model-free:** estimate expectation directly from samples

$$x_i \sim P(x) \qquad E[f(x)] \approx \frac{1}{k} \sum_i f(x_i)$$

- Why does this work? Because samples appear with the right frequencies!

# Model-based Learning

---

- Idea:
  - Learn the model empirically (rather than the “values”)
  - Solve the MDP as if the learned model were correct
  - Better than direct estimation?
- Empirical model learning:
  - Count outcomes for each  $(s,a)$
  - Normalize to give estimate of  $T(s,a,s')$
  - Discover  $R(s,a,s')$  the first time we experience  $(s,a,s')$

# Example: Model-Based Learning

## ■ Episodes:

(1,1) up -1

(1,2) up -1

(1,2) up -1

(1,3) right -1

(2,3) right -1

(3,3) right -1

(3,2) up -1

(3,3) right -1

(4,3) exit +100

(done)

(1,1) up -1

(1,2) up -1

(1,3) right -1


(2,3) right -1

(3,3) right -1

(3,2) up -1

(4,2) exit -100

(done)



$\gamma = 1$

$$T(<3,3>, \text{right}, <4,3>) = 1 / 3$$

$$T(<2,3>, \text{right}, <3,3>) = 2 / 2$$

# Outline

---


- Expectimax Search
- Reinforcement Learning (RL)
- Passive Learning in RL
  - Model-based
  - Model-free
    - Direct Estimation
    - Temporal Difference
- Active Learning in RL
  - Q-Learning



# Model-free Learning

$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$

- **Big idea:** Why bother learning  $T$ ?
- **Question:** How can we compute  $V$  if we don't know  $T$ ?
  - Use direct estimation to sample complete trials
  - Compute “values” for each trial based on the sequence of rewards
  - Average “values” across trials at the end
  - i.e. **sampling!**



# Simple Case: Direct Estimation

## ■ Episodes:

(1,1) up -1

(1,2) up -1

(1,2) up -1

(1,3) right -1

(2,3) right -1

(3,3) right -1

(3,2) up -1

(3,3) right -1

(4,3) exit +100

(done)

(1,1) up -1

(1,2) up -1

(1,3) right -1


(2,3) right -1

(3,3) right -1

(3,2) up -1

(4,2) exit -100

(done)



$\gamma = 1, R = -1$

$$V(2,3) \sim (96 + -103) / 2 = -3.5$$

$$V(3,3) \sim (99 + 97 + -102) / 3 = 31.3$$

# Outline

---


- Expectimax Search
- Reinforcement Learning (RL)
- Passive Learning in RL
  - Model-based
  - Model-free
    - Direct Estimation
    - Temporal Difference
- Active Learning in RL
  - Q-Learning

# Towards Better Model-free Learning

---

## Review: Model-Based Policy Evaluation

- Simplified Bellman updates to calculate  $V$  for a fixed policy:
  - New  $V$  is expected one-step-look-ahead using current  $V$
  - Unfortunately, need  $T$  and  $R$



$$V_0^\pi(s) = 0$$

$$V_{i+1}^\pi(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_i^\pi(s')]$$

# Sample-Based Policy Evaluation?

$$V_{i+1}^{\pi}(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_i^{\pi}(s')]$$

- Who needs T and R? Approximate the expectation with samples (drawn from T!)


$$sample_1 = R(s, \pi(s), s'_1) + \gamma V_i^{\pi}(s'_1)$$

$$sample_2 = R(s, \pi(s), s'_2) + \gamma V_i^{\pi}(s'_2)$$

...

$$sample_k = R(s, \pi(s), s'_k) + \gamma V_i^{\pi}(s'_k)$$


$$V_{i+1}^{\pi}(s) \leftarrow \frac{1}{k} \sum_i sample_i$$



# Model-Difference Learning

---

- **Big idea:** learn from every experience!
  - Update  $V(s)$  each time we experience  $(s, a, s', r)$
  - Likely  $s'$  will contribute updates more often
- Temporal difference learning
  - Policy still fixed!
  - Move values toward value of whatever successor occurs!



**Sample of  $V(s)$ :**  $sample = R(s, \pi(s), s') + \gamma V^\pi(s')$

**Update to  $V(s)$ :**  $V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + (\alpha)sample$

**Same update:**  $V^\pi(s) \leftarrow V^\pi(s) + \alpha(sample - V^\pi(s))$

# Example: TD Policy Evaluation

$$V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + \alpha [R(s, \pi(s), s') + \gamma V^\pi(s')]$$

(1,1) up -1

(1,2) up -1

(1,2) up -1

(1,3) right -1

(2,3) right -1

(3,3) right -1

(3,2) up -1

(3,3) right -1

(4,3) exit +100

(done)

(1,1) up -1

(1,2) up -1

(1,3) right -1


(2,3) right -1

(3,3) right -1

(3,2) up -1

(4,2) exit -100

(done)



Updates for  $V(<3,3>)$ :

$$V(<3,3>) = 0.5 \cdot 0 + 0.5 \cdot [-1 + 1 \cdot 0] = -0.5$$

$$V(<3,3>) = 0.5 \cdot -0.5 + 0.5 \cdot [-1 + 1 \cdot 100] = 49.475$$


$$V(<3,3>) = 0.5 \cdot 49.475 + 0.5 \cdot [-1 + 1 \cdot -0.75]$$

Take  $\gamma = 1$ ,  $\alpha = 0.5$ ,  $V_0(<4,3>) = 100$ ,  $V_0(<4,2>) = -100$ ,  $V_0 = 0$  otherwise

# Problems with TD Value Learning

---

- TD value learning is a model-free way to do policy evaluation
- However, if we want to turn values into a (new) policy, we're sunk:



$$\pi(s) = \arg \max_a Q^*(s, a)$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- **Idea:** learn Q-values directly
- Makes action selection model-free too!



# Outline

---


- Expectimax Search
- Reinforcement Learning (RL)
- Passive Learning in RL
  - Model-based
  - Model-free
    - Direct Estimation
    - Temporal Difference
- Active Learning in RL
  - Q-Learning

# Active Learning

---

- Full reinforcement learning

- You don't know the transitions  $T(s,a,s')$
- You don't know the rewards  $R(s,a,s')$
- You can choose any actions you like
- **Goal: learn the optimal policy**
- ... what value iteration did!



- In this case:

- Learner makes choices!
- Fundamental tradeoff: exploration vs. exploitation
- This is NOT offline planning! You actually take actions in the world and find out what happens...

# Q-Learning Update

---

- Q-Learning: sample-based Q-value iteration

$$Q^*(s, a) = \sum_{s'} T(s, a, s') \left[ R(s, a, s') + \gamma \max_{a'} Q^*(s', a') \right]$$

- Learn  $Q^*(s, a)$  values

- Receive a sample  $(s, a, s', r)$
- Consider your old estimate:  $Q(s, a)$
- Consider your new sample estimate:

$$sample = R(s, a, s') + \gamma \max_{a'} Q(s', a')$$

- Incorporate the new estimate into a running average:

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + (\alpha) [sample]$$

# Exploration / Exploitation



---

- Several schemes for forcing exploration
  - Simplest: random actions ( $\epsilon$  greedy)
    - Every time step, flip a coin
    - With probability  $\epsilon$ , act randomly
    - With probability  $1-\epsilon$ , act according to current policy
  - Problems with random actions?
    - You do explore the space, but keep thrashing around once learning is done
    - One solution: lower  $\epsilon$  over time
    - Another solution: exploration functions

# Q-Learning Properties

---

- Amazing result: Q-learning converges to optimal policy
  - If you explore enough
  - If you make the learning rate small enough
  - ... but not decrease it too quickly!
  - Basically doesn't matter how you select actions (!)
- Neat property: off-policy learning
  - learn optimal policy without following it (some caveats)



# RL for Helicopter Controller

---

